

DIVYANSHU GOYAL

Sunnyvale, CA ♦ (+1) 404-414-3762 ♦ divyanshu.g25@gmail.com

[LinkedIn](#) ♦ [GitHub](#) ♦ [Google Scholar](#)

SUMMARY

Applied Scientist and multimodal researcher with 10 years at the intersection of frontier ML research and production engineering. I lead the training and fine-tuning of **vision-language models** at scale, design novel **evaluation and alignment** methods (**5 patents**, peer-reviewed and arXiv publications), and previously led the AI powering **Adobe GenStudio for Performance Marketing**; now serving Fortune 500 clients including Microsoft and AT&T. Deep expertise across **VLM training**, **multimodal reasoning**, and **LLM pretraining & scaling laws**, paired with the systems engineering to deploy them in production.

EXPERIENCE

Applied Scientist – Adobe Unified Platform, San Jose

June 2021 – Present

♦ *Vision-Language Model Training & Evaluation*

- Fine-tuned a **Gemma-3-12B vision-language model** with **LoRA** across **8× H100 GPUs** (DDP, bfloat16, Flash Attention 2 + Liger kernels) for large-scale marketing-quality assessment.
- Pioneered an **observable injection** technique that grounds VLM quality reasoning in **14+ computed pixel-level metrics** (sharpness, exposure, color balance); built an interpretable scorer over these signals reaching **Pearson $r \approx 0.88$** at **MAE 0.77**, outperforming GPT-5.4 (MAE 1.34) and Claude Opus 4.7 (MAE 1.95) baselines.
- Developed the **MQCore evaluation benchmark** (inspired by DCLM CORE) and a fault-tolerant async annotation pipeline processing **100+ images/min** over the **700K-image KADIS** corpus.

♦ *Multimodal Research & Benchmarking*

- Authored **DistortBench** (CVPR FGVC13 Workshop 2026): a **13.5K-question** diagnostic benchmark spanning **27 distortion types**, 6 perceptual categories, and 5 severity levels.
- Evaluated **18+ frontier VLMs** (GPT-5.4, Claude Sonnet 4.6, Qwen, InternVL, Gemma) with bootstrap confidence intervals and balanced-accuracy metrics, exposing low-level perception as a major weakness: top model **61.9%** vs. **65.7% human baseline**.
- Created a curiosity-driven **LLM-as-a-judge** framework for personalized creative evaluation using Bayesian surprise as a quantifiable novelty signal (arXiv 2025), achieving **23% F1** and **38% correlation** improvements over baselines.

♦ *Production Multimodal Systems & Agents*

- Architected the **Unified Brand Service** (90K+ LOC across 3 global regions) powering **Adobe GenStudio**: multimodal brand-DNA extraction, RAG-based retrieval, and GPT-4-Vision content validation.
- Built a **multimodal agent** that drives Adobe Photoshop from natural language via a **49-tool** observe-think-act loop with vision grounding and an iterative quality-critic refinement layer.
- Built the founding prototype of **Adobe GenStudio for Performance Marketing** (launched Oct 2024), now serving Fortune 500 clients including Microsoft and AT&T; invented patented VLM pipelines for **cross-cultural image adaptation** and **constrained machine translation** (US Patents 2026).

Machine Learning Engineer – Adobe Inc., Bangalore

July 2016 – August 2019

- Co-engineered Adobe's production **ML inference platform**, sustaining **3,500 QPS/node** at **0.3–0.6ms p99** latency through flame-graph profiling, JMH micro-benchmarking, and lock-free design.
- Architected a **GC-free asynchronous event pipeline** on the LMAX Disruptor ring buffer, decoupling telemetry from prediction to eliminate back-pressure at peak load.
- Built an end-to-end **model lifecycle platform**: versioned deployment, live monitoring, and staged rollout across dev, staging, and production.
- Co-authored the [Adobe Tech Blog](#) on engineering high-throughput, low-latency ML services.

PUBLICATIONS & PATENTS

Curiosity-Driven LLM-as-a-judge for Personalized Creative Judgment
V. Bannihatti Kumar, **D. Goyal**, A. Eppa, N. Bhandari
A curiosity-driven framework that personalizes creative-writing evaluation to individual judgments via LLM-as-a-judge, targeting the subjectivity gap in LLM assessment.

arXiv 2025
[\[paper\]](#)

Cultural Adaptation of Images Using Generative AI
A. Eppa, M. Anand, **D. Goyal**
Multi-stage pipeline where a vision-language model grounds source images as text, conditioned on region-specific cultural guidelines, and a text-to-image model renders culturally adapted output.

US Patent App. 2026
[\[patent\]](#)

Guiding Language Translation with Translation Documents Using ML
D. Goyal, A. Eppa, M. Anand
Structured approach to constrained machine translation by extracting guidelines into language-agnostic facets and computing facet-specific adherence scores for linguistic and stylistic alignment.

US Patent App. 2026
[\[patent\]](#)

Parameter-Efficient Model Serving Systems
Y. Wang, N. Vangala, M. Anand, **D. Goyal**, et al.
System for efficiently serving fine-tuned models by dynamically loading parameter-efficient layers into pre-loaded base models, minimizing network transfer costs and container startup time.

US Patent App. 2025
[\[patent\]](#)

Brand-Aligned Content Generation with Structured Knowledge
M. Anand, J. Mathew, **D. Goyal**
Generative system that transforms unstructured brand information into structured brand DNA elements with confidence scoring, enabling brand-aligned content creation across channels.

US Patent App. 2025
[\[patent\]](#)

SELECTED PROJECTS & TECHNICAL WRITING

VibeNanoChat:GPT-2 Scale LLM Pretraining Framework
From-scratch decoder-only transformer (RoPE, RMSNorm, sliding-window attention) with one-knob depth-parameterized scaling, a custom distributed **Muon** optimizer with ZeRO-2 state sharding, and Flash-Attention-3 / bfloat16 training on 8x H100. Empirically verified the **C = 6ND** compute law across 40+ scaling-law runs.

2026
[\[code\]](#)

A Lazy Engineer’s Guide to Scaling Transformers
Deriving all transformer architecture dimensions, training hyperparameters, and compute schedules from a single integer (depth), enabling trivial multi-scale training sweeps.

March 2026
[\[blog\]](#)

Going Fast: Every Optimization That Made LLM Training Fly
Six concrete optimizations:TF32, BF16 mixed precision, torch.compile, Flash Attention, parallel DataLoaders, and pinned memory:that collectively drove significant MFU improvements during GPT-2 scale training.

March 2026
[\[blog\]](#)

TECHNICAL SKILLS

Research:	LLM Pretraining & Alignment (SFT, RLHF, DPO), Vision-Language Models, Multimodal Agents, Evaluation & Benchmarking, Data Curation
Training & Serving:	PyTorch, Hugging Face Transformers/PEFT, DeepSpeed, FSDP, vLLM, Flash Attention, LoRA/QLoRA, CUDA, Weights & Biases
Infrastructure:	Docker, LangChain, RAG pipelines
Programming:	Python, C++, JavaScript

EDUCATION

Georgia Institute of Technology – M.S. Computer Science (Machine Learning) 2019 – 2021 | **GPA: 4.0**
BITS Pilani – B.E. Computer Science, M.Sc. Mathematics 2011 – 2016 | **GPA: 8.75/10**