DIVYANSHU GOYAL

Sunnyvale, CA ◇ +1 404-414-3762 ◇ divyanshu.g25@gmail.com
linkedin.com/in/divyanshu2594 ◇ github.com/divyanshu25 ◇ Google Scholar

SUMMARY

I am an Applied Scientist at Adobe, where I train and evaluate multimodal language models for creative work. Much of it deals with problems where quality is subjective, such as whether a photo edit looks better or a story feels original, so I spend as much time building evaluations I can trust as training the models themselves. I have shipped this work in Adobe GenStudio and published it at CVPR, ICML, and AAAI workshops.

EXPERIENCE

Applied Scientist – Adobe Inc. (Unified Platform), San Jose, CA June 2021 – Present

Vision-Language Model Training & Evaluation

- Led **ATune** (under review): a 54M-parameter model that learns photo retouching from 10M unedited stock photographs and drives off-the-shelf editors through their sliders. In a study with 157 raters, its edits were preferred over GPT-6 Astra's in **54.3%** of votes (95% CI 50.3–58.3) and tied with Claude Opus 5's (50.5%, CI 46.8–54.2).
- Fine-tuned **Gemma-3-12B** with LoRA on 8× H100 (DDP, BF16, FlashAttention-2, Liger kernels) to assess the quality of marketing images.
- Grounded VLM quality reasoning in 14+ computed pixel-level metrics (sharpness, exposure, color balance); the resulting interpretable scorer reaches **Pearson r ≈ 0.88** (MAE 0.77) vs. zero-shot GPT-5.4 (MAE 1.34) and Claude Opus 4.7 (MAE 1.95).
- Built the **MQCore** evaluation benchmark and a fault-tolerant async annotation pipeline (100+ images/min) for the 700K-image KADIS image-quality dataset.

Evaluation Research & Benchmarking

- First author of **DistortBench** (CVPR 2026 Workshop): a 13.5K-question benchmark over 27 distortion types and 5 severity levels. Evaluated 18 VLMs with bootstrap confidence intervals; the best model (Qwen3.5-27B) scores **61.9%** (chance 25%), above the average individual annotator (60.2%) but below a 3-annotator majority vote (65.7%).
- Co-authored **JudgeBreak** (ICML 2026 Workshop): across 95 subjective evaluation settings, linear probes on Qwen3-32B hidden states with cross-validated layer selection reach **78.1%** mean pairwise accuracy vs. 70.1% for the strongest baseline (the best reward model or a zero-shot GPT-5.4 judge).
- Co-authored a curiosity-driven **LLM-as-a-judge** that personalizes creative-writing evaluation using an intrinsic curiosity model of the judge's surprise at each expert's explanation (AAAI 2026 Workshop, Oral), improving F1 by 23% and correlation by 38% over an SFT baseline.

Production Multimodal Systems & Agents

- Built the founding prototype of **Adobe GenStudio for Performance Marketing** (launched October 2024, now serving Fortune 500 clients including Microsoft and AT&T) and architected its multimodal brand service (brand-DNA extraction, retrieval, VLM content validation) across 3 regions.
- Built a **multimodal Photoshop agent** that edits images from natural language through a 49-tool observe, think, and act loop with vision grounding and a quality-critic refinement step.

Machine Learning Engineer – Adobe Inc., Bangalore, India July 2016 – August 2019

- Co-engineered Adobe's production **ML inference platform**, sustaining 3,500 QPS/node at 0.3–0.6 ms app-side p99 latency (Adobe Tech Blog).
- Built an end-to-end model lifecycle platform: versioned deployment, live monitoring, and staged rollout across dev, staging, and production.

SELECTED PROJECTS

VibeNanoChat: GPT-2 Scale LLM Pretraining from Scratch 2026 | [code]
Decoder-only transformer (RoPE, RMSNorm, sliding-window attention) with one-knob depth scaling, a custom distributed **Muon** optimizer with ZeRO-2 state sharding, and FlashAttention-3 and BF16 training on 8× H100 at over 1M tokens/sec. Fit isoFLOP curves across 40+ runs (44M to 522M parameters) to find compute-optimal model sizes.
Write-ups: Going Fast: Every Optimization That Made LLM Training Fly and A Lazy Engineer's Guide to Scaling Transformers.

PUBLICATIONS

DistortBench: Benchmarking Vision Language Models on Image Distortion Identification
D. Goyal, A. Eppa, V. Bannihatti Kumar CVPR 2026 Workshop | [paper]

JudgeBreak: Is Taste Hidden in the Layers? Stress-Testing Reward Models and LLM Judges
V. Bannihatti Kumar, A. Narang, M. Anand, A. Eppa, **D. Goyal** ICML 2026 Workshop | [paper]

Curiosity-Driven LLM-as-a-judge for Personalized Creative Judgment
V. Bannihatti Kumar, **D. Goyal**, A. Eppa, N. Bhandari AAAI 2026 Workshop (Oral) | [paper]

PATENTS

Granted: Guiding Language Translation with Translation Documents Using ML (first inventor, 2026) [patent]

Pending: Cultural Adaptation of Images Using Generative AI (2026) [patent]; Parameter-Efficient Model Serving (2025) [patent]; Brand-Aligned Content Generation (2025) [patent]

TECHNICAL SKILLS

Research: VLM Training & Evaluation, LLM Pretraining & Scaling Laws, Benchmark Design, Human Preference Studies, LLM-as-a-judge & Reward Models, Linear Probing, Bootstrap Statistics

Training: Python, PyTorch, Distributed Training (DDP, ZeRO-2), torch.compile, FlashAttention-2/3, Liger kernels, Muon, Mixed Precision (BF16, TF32), LoRA, Hugging Face Transformers/PEFT, Weights & Biases

EDUCATION

Georgia Institute of Technology – M.S. Computer Science (Machine Learning), GPA 4.0 (2019 – 2021)

BITS Pilani – B.E. Computer Science, M.Sc. Mathematics, GPA 8.75/10 (2011 – 2016)